



Open in app

Get started



Published in Towards Data Science

You have **1** free member-only story left this month. [Sign up for Medium and get an extra one](#)



Serafeim Loukas

Follow

May 30, 2020 · 7 min read · ✨ · 🎧 Listen



Save



# PCA clearly explained —When, Why, How to use it and feature importance: A guide in Python

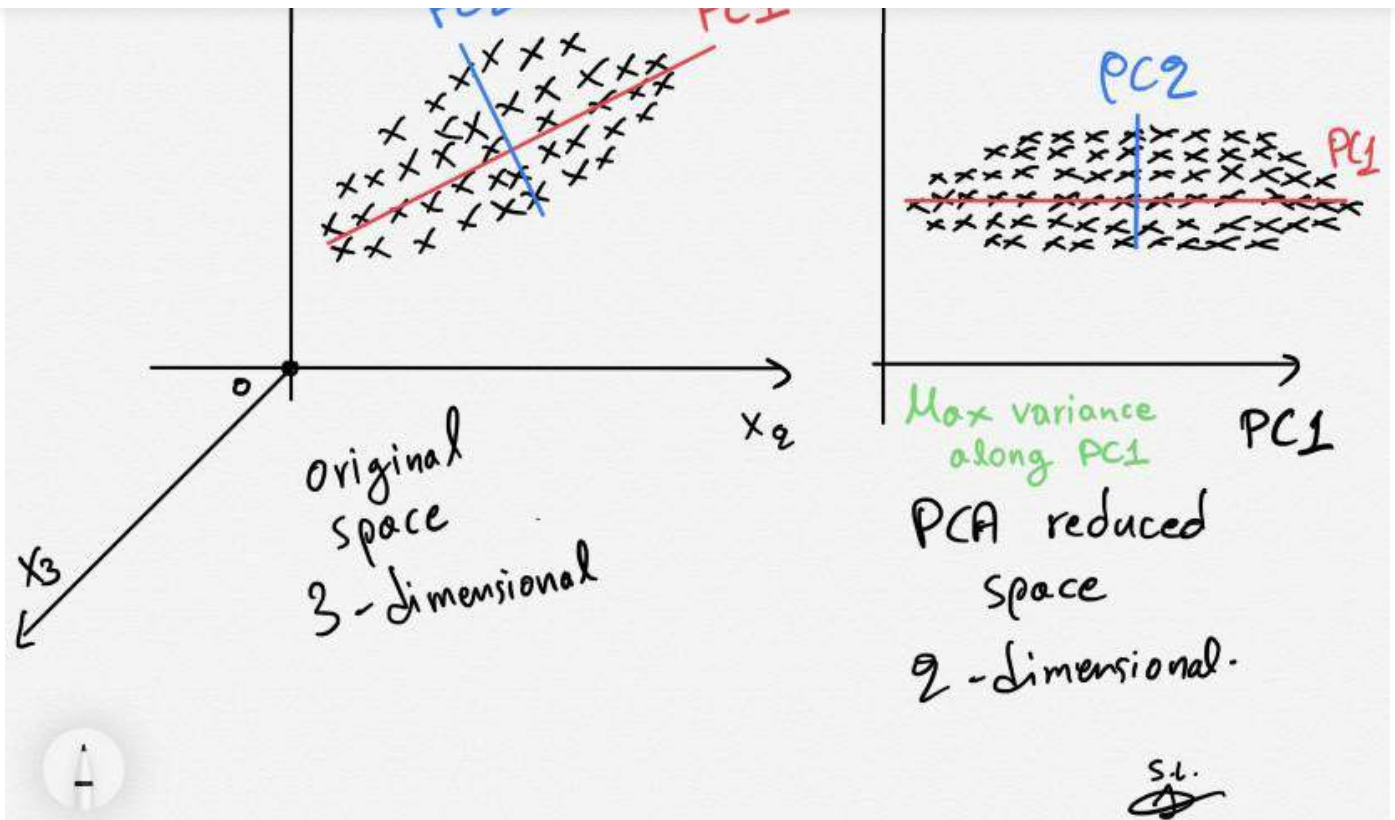
In this post I explain what PCA is, when and why to use it and how to implement it in Python using scikit-learn. Also, I explain how to get the feature importance after a PCA analysis.





Open in app

Get started



Handmade sketch made by the author.

## 1. Introduction & Background

Principal Components Analysis (PCA) is a well-known **unsupervised dimensionality reduction** technique that constructs **relevant** features/variables through linear (linear PCA) or non-linear (kernel PCA) **combinations** of the original variables (features). In this post, we will only focus on the famous and widely used **linear PCA** method.

The construction of relevant features is achieved by **linearly transforming correlated variables** into a smaller number of **uncorrelated** variables. This is done by **projecting** (dot product) the original data into the **reduced PCA space** using the eigenvectors of the covariance/correlation matrix aka the principal components (PCs).

The **resulting projected data** are essentially **linear combinations** of the **original data** capturing most of the variance in the data ([Jolliffe 2002](#)).

In summary, **PCA** is an **orthogonal transformation** of the data into a series of





Open in app

Get started

- My mailing list in just 5 seconds: <https://seralouk.medium.com/subscribe>

- Become a member and support me: <https://seralouk.medium.com/membership>

## 2. When/Why to use PCA

- PCA technique is particularly useful in processing data where **multi-collinearity** exists between the **features/variables**.
- PCA can be used when the **dimensions of the input features are high** (e.g. a lot of variables).
- PCA can be also used for **denoising** and **data compression**.

## 3. Core of the PCA method

Let  $\mathbf{x}$  be a matrix containing the original data with shape  $[n\_samples, n\_features]$  .

Briefly, the PCA analysis consists of the following **steps**:

- First, the original input variables stored in  $\mathbf{x}$  are **z-scored** such each original variable (column of  $\mathbf{x}$ ) has zero mean and unit standard deviation.
- The next step involves the construction and **eigendecomposition** of the **covariance** matrix  $\mathbf{C}_\mathbf{x} = (1/n)\mathbf{x}'\mathbf{x}$  (in case of z-scored data the covariance is equal to the correlation matrix since the standard deviation of all features is 1).
- **Eigenvalues** are then **sorted** in a **decreasing** order representing decreasing variance in the data (the eigenvalues are equal to the variance — I will prove this below using Python in Paragraph 6).





Open in app

Get started

normalized data by the **leading eigenvectors** of the covariance matrix i.e. the PCs.

- The new **reduced** PCA space **maximizes** the **variance** of the **original** data. To **visualize** the projected data as well as the contribution of the original variables, in a joint plot, we can use the **biplot**.

#### 4. The maximum number of meaningful components

There is an **upper bound** of the **meaningful components** that can be extracted using PCA. This is related to the **rank** of the **covariance/correlation matrix** ( $C_X$ ). Having a data matrix  $X$  with shape  $[n\_samples, n\_features/n\_variables]$ , the **covariance/correlation matrix** would be  $[n\_features, n\_features]$  with **maximum rank** equal to  $\min(n\_samples, n\_features)$ .

Thus, we can have at **most**  $\min(n\_samples, n\_features)$  **meaningful PC components/dimensions** due to the **maximum rank** of the covariance/correlation matrix.

#### 5. Python example using scikit-learn and the Iris dataset

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.decomposition import PCA
import pandas as pd
from sklearn.preprocessing import StandardScaler
plt.style.use('ggplot')

# Load the data
iris = datasets.load_iris()
X = iris.data
y = iris.target

# Z-score the features
scaler = StandardScaler()
scaler.fit(X)
X = scaler.transform(X)
```





Open in app

Get started

Let's plot the data before and **after** the PCA transform and also **color** code each point (sample) using the corresponding **class of the flower** ( $y$ ) .

```
fig, axes = plt.subplots(1,2)
axes[0].scatter(X[:,0], X[:,1], c=y)
axes[0].set_xlabel('x1')
axes[0].set_ylabel('x2')
axes[0].set_title('Before PCA')

axes[1].scatter(X_new[:,0], X_new[:,1], c=y)
axes[1].set_xlabel('PC1')
axes[1].set_ylabel('PC2')
axes[1].set_title('After PCA')



plt.show()
```



533



7





Open in app

Get started

```
print(pca.explained_variance_ratio_)
# array([0.72962445, 0.22850762])
```

## 6. Proof of eigenvalues of original covariance matrix being equal to the variances of the reduced space

### Mathematical formulation & proof

Assuming that the original input variables stored in  $\mathbf{x}$  are z-scored such each original variable (column of  $\mathbf{x}$ ) has zero mean and unit standard deviation, we have:

- The covariance of the original space is:  $\mathbf{C}_\mathbf{X} = \frac{1}{n} \mathbf{X}^T \mathbf{X}$

- PCA: Eigendecomposition of the covariance of the original space is:

$$\mathbf{C}_\mathbf{X} = \frac{1}{n} \mathbf{X}^T \mathbf{X} = \mathbf{U} \mathbf{\Lambda} \mathbf{U}^T$$

- Let  $\mathbf{Y} = \mathbf{XU}$  be the projected data ( $\mathbf{U}$  stores the eigenvectors as columns)

- The covariance of the reduced PCA space is:

$$\begin{aligned} \mathbf{C}_\mathbf{Y} &= \frac{1}{n} \mathbf{Y}^T \mathbf{Y} \\ &= \frac{1}{n} (\mathbf{XU})(\mathbf{XU})^T \\ &= \frac{1}{n} \mathbf{U}^T \mathbf{X}^T \mathbf{X} \mathbf{U} \\ &= \mathbf{U}^T \left( \frac{1}{n} \mathbf{X}^T \mathbf{X} \right) \mathbf{U} \\ \mathbf{C}_\mathbf{Y} &= \mathbf{U}^T \mathbf{C}_\mathbf{X} \mathbf{U} \end{aligned}$$



[Open in app](#)[Get started](#)

## Verify using Python

The **maximum variance proof** can be also seen by estimating the **covariance** matrix of the **reduced space**:

```
np.cov(X_new.T)

array([[2.93808505e+00, 4.83198016e-16],
       [4.83198016e-16, 9.20164904e-01]])
```

We observe that these values (on the diagonal we have the variances) are **equal** to the **actual eigenvalues** of the covariance stored in `pca.explained_variance_`:

```
pca.explained_variance_
array([2.93808505, 0.9201649 ])
```

## 7. Feature importance

The **importance** of each **feature** is reflected by the **magnitude** of the **corresponding values in the eigenvectors** (higher magnitude — higher importance).

Let's find the **most important features**:

```
print(abs( pca.components_ ))

[[0.52106591 0.26934744 0.5804131 0.56485654]
 [0.37741762 0.92329566 0.02449161 0.06694199]]
```

Here, `pca.components_` has shape `[n_components, n_features]` Thus, by looking at the **PC1** (first Principal Component) which is the **first row**





Open in app

Get started

can state that **feature 2** and then **1** are the **most important** for **PC2**.

*To sum up, we look at the absolute values of the eigenvectors' components corresponding to the  $k$  largest eigenvalues. In sklearn the components are sorted by explained variance. The larger they are these absolute values, the more a specific feature contributes to that principal component.*

## 8. The biplot

The **biplot** is the best way to visualize **all-in-one** following a **PCA** analysis.

There is an implementation in **R** but there is no standard implementation in **python** so I decided to write my **own function** for that:

```
def biplot(score, coeff , y):
    '''
    Author: Serafeim Loukas, serafeim.loukas@epfl.ch
    Inputs:
        score: the projected data
        coeff: the eigenvectors (PCs)
        y: the class labels
    '''

    xs = score[:,0] # projection on PC1
    ys = score[:,1] # projection on PC2
    n = coeff.shape[0] # number of variables
    plt.figure(figsize=(10,8), dpi=100)
    classes = np.unique(y)
    colors = ['g','r','y']
    markers=['o','^','x']
    for s,l in enumerate(classes):
        plt.scatter(xs[y==l],ys[y==l], c = colors[s],
marker=markers[s]) # color based on group
    for i in range(n):
        #plot as arrows the variable scores (each variable has a score
for PC1 and one for PC2)
        plt.arrow(0, 0, coeff[i,0], coeff[i,1], color = 'k', alpha =
0.9,linestyle = '-',linewidth = 1.5, overhang=0.2)
        plt.text(coeff[i,0]* 1.15, coeff[i,1] * 1.15, "Var"+str(i+1),
color = 'k', ha = 'center', va = 'center',fontsize=10)
```





[Open in app](#)[Get started](#)

```
plt.ylim([-limy, limy])  
plt.grid()  
plt.tick_params(axis='both', which='both', labelsize=14)
```

Call the function (make sure to run first the initial blocks of code where we load the iris data and perform the PCA analysis):


```
import matplotlib as mpl  
mpl.rcParams.update(mpl.rcParamsDefault) # reset ggplot style  
  
# Call the biplot function for only the first 2 PCs  
biplot(X_new[:,0:2], np.transpose(pca.components_[0:2, :]), y)  
plt.show()
```





Open in app

Get started



The PCA biplot using my custom function.

We can again verify **visually** that **a)** the variance is maximized and **b)** that **feature 1, 3 and 4** are the **most important** for **PC1**. Similarly, **feature 2** and then **1** are the **most important** for **PC2**.

Furthermore, **arrows** (variables/features) that point into the **same direction** indicate **correlation** between the variables that they represent whereas, the arrows heading in **opposite directions** indicate a **contrast** between the variables they represent.

Verify the above using **code**:

```
# Var 3 and Var 4 are extremely positively correlated
np.corrcoef(X[:, 2], X[:, 3])[1, 0]
```



[Open in app](#)[Get started](#)

That's all folks! Hope you liked this article!

## Latest posts

### Time-Series Forecasting: Predicting Stock Prices Using Facebook's Prophet Model

Predict stock prices using a forecasting model publicly available from Facebook: The Prophet

[towardsdatascience.com](#)

### ROC Curve Explained using a COVID-19 hypothetical example: Binary & Multi-Class Classification...

In this post I clearly explain what a ROC curve is and how to read it. I use a COVID-19 example to make my point and I...

[towardsdatascience.com](#)

### Support Vector Machines (SVM) clearly explained: A python tutorial for classification problems...

In this article I explain the core of the SVMs, why and how to use them. Additionally, I show how to plot the support...

[towardsdatascience.com](#)

### Everything you need to know about Min-Max normalization in Python

In this post I explain what Min-Max scaling is, when to use it and how to





Open in app

Get started

## How Scikit-Learn's StandardScaler works

In this post I am explaining why and how to apply Standardization using scikit-learn

towardsdatascience.com

## Stay tuned & support this effort

If you liked and found this article useful, **follow** me!

Questions? Post them as a comment and I will reply as soon as possible.

## References

[1] Jolliffe, I. T. *Principal component analysis*. New York, NY: Springer, 2002.

[2] [https://en.wikipedia.org/wiki/Principal\\_component\\_analysis](https://en.wikipedia.org/wiki/Principal_component_analysis)

[3] <https://stattrek.com/matrix-algebra/matrix-rank.aspx>

## Get in touch with me

- LinkedIn: <https://www.linkedin.com/in/serafeim-loukas/>
- ResearchGate: [https://www.researchgate.net/profile/Serafeim\\_Loukas](https://www.researchgate.net/profile/Serafeim_Loukas)
- EPFL profile: <https://people.epfl.ch/serafeim.loukas>
- Stack Overflow: <https://stackoverflow.com/users/5025009/seralouk>

